

09 - Classification

Logistic Regression, Discriminant Analysis, and Naive Bayes

SYS 6018 | Fall 2019

09-classification.pdf

Contents

1	Classification Intro	2
1.1	Credit Card Default data (Default)	2
2	Classification and Pattern Recognition	5
2.1	Binary Classification	5
3	Logistic Regression	8
3.1	Basics	8
3.2	Estimation	8
3.3	Logistic Regression in Action	9
4	Linear/Quadratic Discriminant Analysis (LDA/QDA)	11
4.1	Estimation	11
4.2	LDA/QDA in Action	12
5	Evaluating Classification Models	12
5.1	Evaluation of Binary Classification Models	13
5.2	Common Binary Loss Functions	14
5.3	Evaluating Binary Classification Models	15
5.4	Performance Metrics	16
5.5	Performance over a range of thresholds	16
6	Naive Bayes and Generalized Additive Models (GAM)	19
6.1	The Bayes Breakdown	19
6.2	Naive Bayes	20
6.3	Logistic Regression vs. Linear Discriminant Analysis (LDA) vs. Naive Bayes	21

Some of the figures in this presentation are taken from "An Introduction to Statistical Learning, with applications in R" (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

1 Classification Intro

1.1 Credit Card Default data (Default)

The textbook *An Introduction to Statistical Learning (ISL)* has a description of a simulated credit card default dataset. The interest is on predicting whether an individual will default on their credit card payment.

```
data(Default, package="ISLR")
```

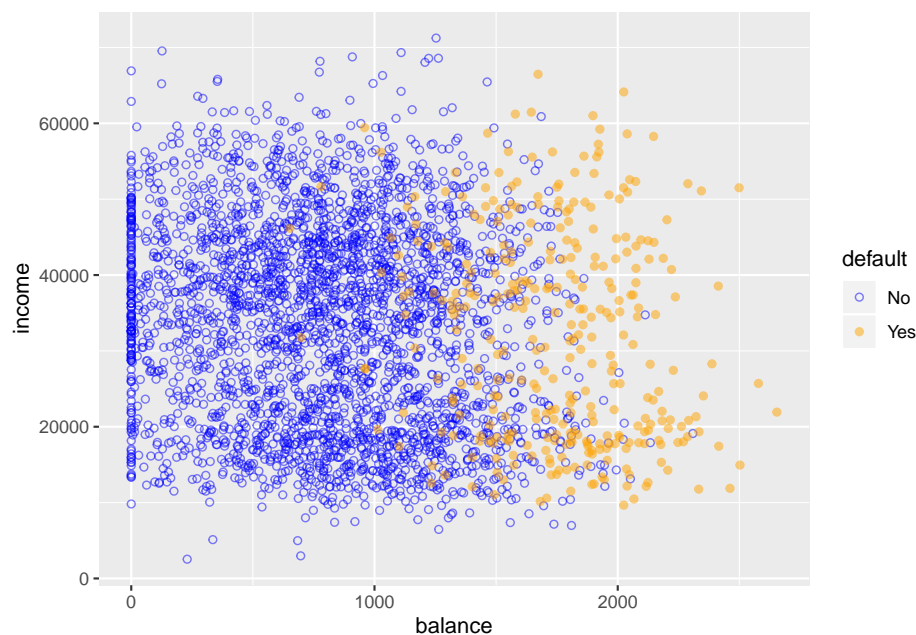
The variables are:

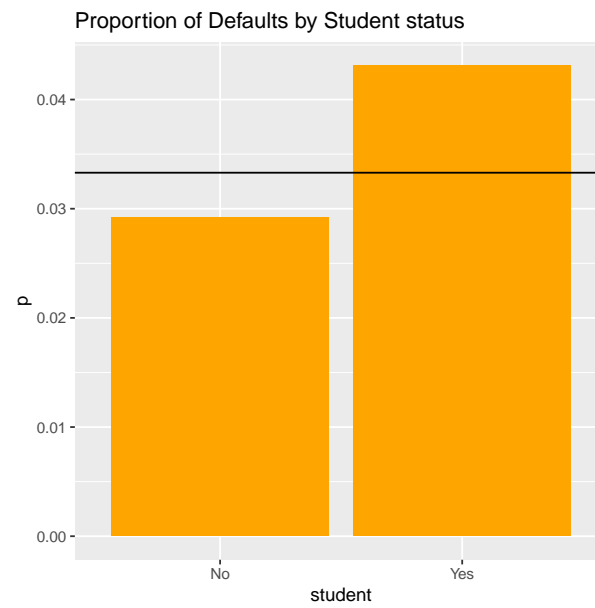
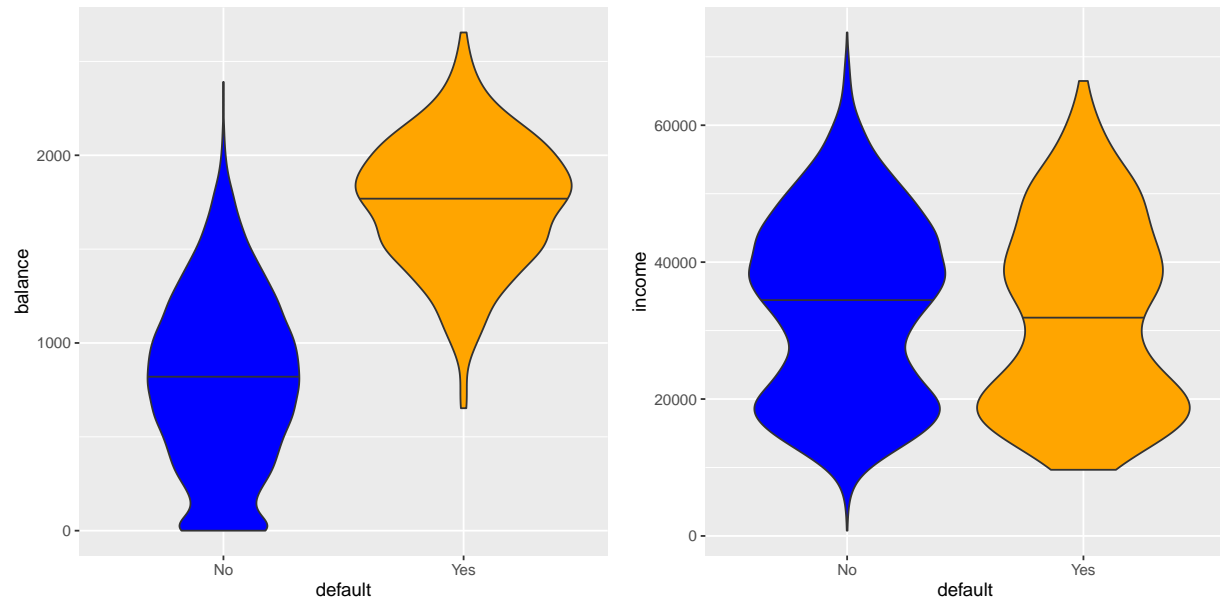
- *response variable* is categorical (factor) Yes and No, (default)
- the categorical (factor) variable (*student*) is either Yes or No
- the average balance a customer has after making their monthly payment (*balance*)
- the customer's income (*income*)

default	student	balance	income
No	No	729.5	44362
No	Yes	817.2	12106
No	No	1073.5	31767
No	No	529.3	35704
No	No	785.7	38463
No	Yes	919.6	7492

```
summary(Default)
```

```
#> default student balance income
#> No :9667 No :7056 Min. : 0 Min. : 772
#> Yes: 333 Yes:2944 1st Qu.: 482 1st Qu.:21340
#> Median : 824 Median :34553
#> Mean : 835 Mean :33517
#> 3rd Qu.:1166 3rd Qu.:43808
#> Max. :2654 Max. :73554
```





Your Turn #1 : Credit Card Default Modeling

How would you construct a model to predict defaults?

2 Classification and Pattern Recognition

- The response variable is categorical and denoted $G \in \mathcal{G}$
 - Default Credit Card Example: $\mathcal{G} = \{\text{"Yes"}, \text{"No"}\}$
 - Medical Diagnosis Example: $\mathcal{G} = \{\text{"stroke"}, \text{"heart attack"}, \text{"drug overdose"}, \text{"vertigo"}\}$
- The training data is $D = \{(X_1, G_1), (X_2, G_2), \dots, (X_n, G_n)\}$
- The optimal decision/classification is often based on the posterior probability $\Pr(G = g \mid \mathbf{X} = \mathbf{x})$

2.1 Binary Classification

- Classification is simplified when there are only 2 classes.
 - Many multi-class problems can be addressed by solving a set of binary classification problems (e.g., [one-vs-rest](#)).
- It is often convenient to *code* the response variable to a binary $\{0, 1\}$ variable:

$$Y_i = \begin{cases} 1 & G_i = \mathcal{G}_1 \quad (\text{outcome of interest}) \\ 0 & G_i = \mathcal{G}_2 \end{cases}$$

- In the `Default` data, it would be natural to set `default=Yes` to 1 and `default=No` to 0.

2.1.1 Linear Regression

- In this set-up we can run linear regression

$$\hat{y}(\mathbf{x}) = \hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j x_j$$

```
#-- Create binary column (y)
Default = Default %>% mutate(y = ifelse(default == "Yes", 1L, 0L))

#-- Fit Linear Regression Model
fit.lm = lm(y ~ student + balance + income, data=Default)
```

term	estimate	std.error	statistic	p.value
(Intercept)	-0.0812	0.0084	-9.685	0.0000
studentYes	-0.0103	0.0057	-1.824	0.0682
balance	0.0001	0.0000	37.412	0.0000
income	0.0000	0.0000	1.039	0.2990

Your Turn #2 : OLS for Binary Responses

1. For the binary Y , what is linear regression estimating?

2. What is the *loss function* that linear regression is using?
3. How could you create a *hard classification* from the linear model?
4. Does it make sense to use linear regression for binary classification?

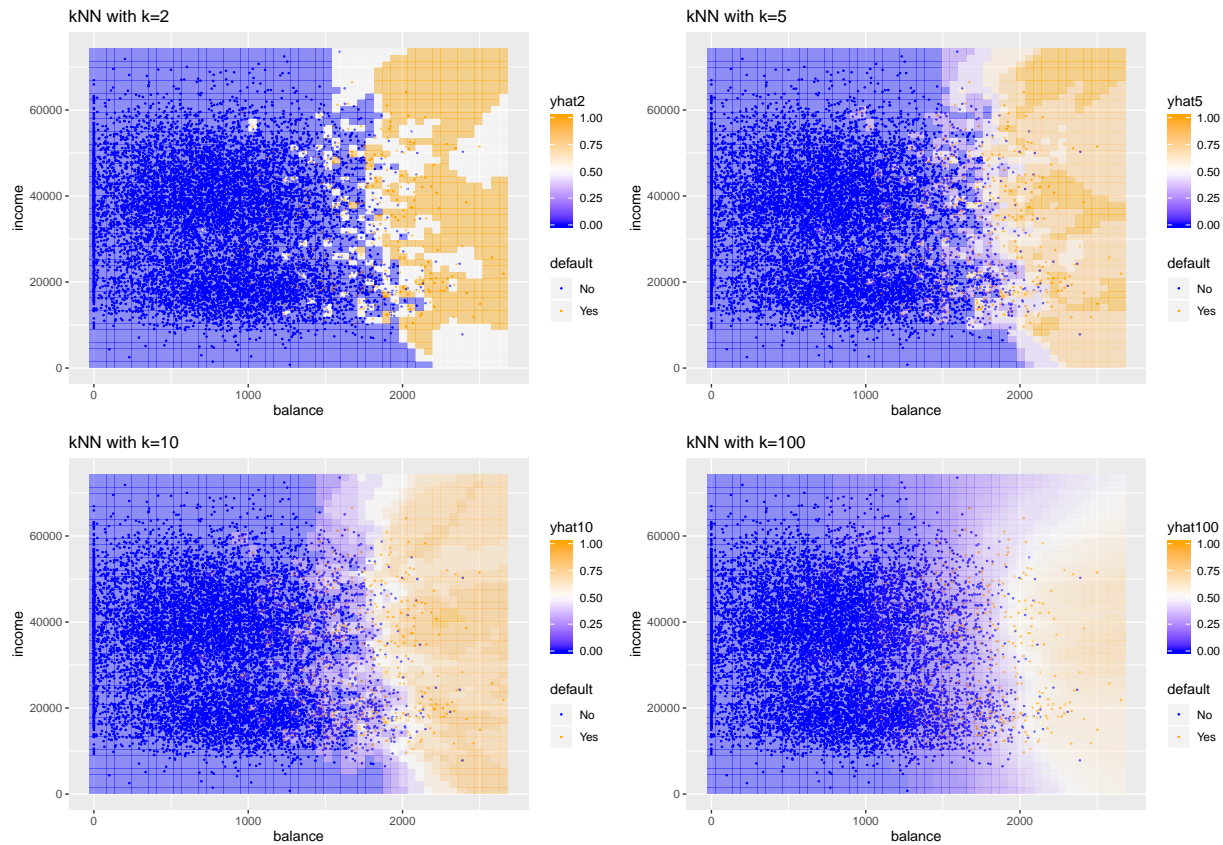
2.1.2 *k*-nearest neighbor (kNN)

- The *k*-NN method is a non-parametric *local* method, meaning that to make a prediction $\hat{y}|x$, it only uses the training data in the *vicinity* of x .
 - contrast with OLS linear regression, which uses all X 's to get prediction.
- The model (for regression and binary classification) is simple to describe

$$\begin{aligned} f_{\text{knn}}(x; k) &= \frac{1}{k} \sum_{i: x_i \in N_k(x)} y_i \\ &= \text{Avg}(y_i \mid x_i \in N_k(x)) \end{aligned}$$

- $N_k(x)$ are the set of k nearest neighbors
 - only the k closest y 's are used to generate a prediction
 - it is a *simple mean* of the k nearest observations
- When y is binary (i.e., $y \in \{0, 1\}$), the kNN model estimates

$$f_{\text{knn}}(x; k) \approx p(x) = \Pr(Y = 1 | X = x)$$



Your Turn #3 : Thoughts about kNN

The above plots show a kNN model using the *continuous* predictors of *balance* and *income*.

- How could you use kNN with the categorical *student* predictor?

- The k -NN model also has a more general description when the response variables is categorical $G_i \in \mathcal{G}$

$$f_g^{\text{knn}}(x; k) = \frac{1}{k} \sum_{i: x_i \in N_k(x)} \mathbb{1}(g_i = g) \\ = \widehat{\Pr}(G_i = g \mid x_i \in N_k(x))$$

- $N_k(x)$ are the set of k nearest neighbors
- only the k closest y 's are used to generate a prediction
- it is a *simple proportion* of the k nearest observations that are of class g

3 Logistic Regression

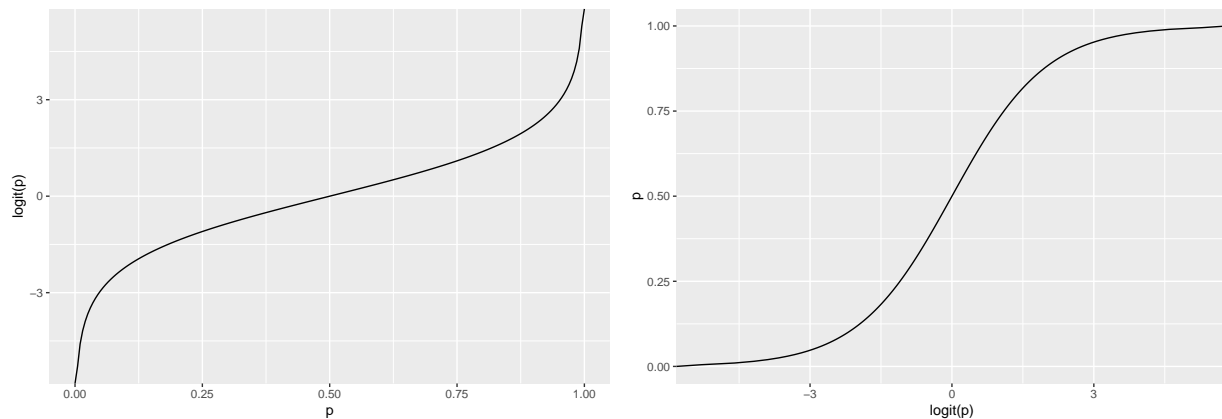
3.1 Basics

- Let $0 \leq p \leq 1$ be a probability.
- The log-odds of p is called the *logit*

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

- The inverse logit is the *logistic function*. Let $f = \text{logit}(p)$, then

$$\begin{aligned} p &= \frac{e^f}{1 + e^f} \\ &= \frac{1}{1 + e^{-f}} \end{aligned}$$



- For binary response variables $Y \in \{0, 1\}$, linear regression models estimate

$$E[Y \mid X = x] = \Pr(Y = 1 \mid X = x) = \beta^\top x$$

- Logistic Regression models alternatively estimate

$$\log\left(\frac{\Pr(Y = 1 \mid X = x)}{1 - \Pr(Y = 1 \mid X = x)}\right) = \beta^\top x$$

and thus,

$$\begin{aligned} \Pr(Y = 1 \mid X = x) &= \frac{e^{\beta^\top x}}{1 + e^{\beta^\top x}} \\ &= \left(1 + e^{-\beta^\top x}\right)^{-1} \end{aligned}$$

3.2 Estimation

- The data for logistic regression is: $(\mathbf{x}_i, y_i)_{i=1}^n$ where $y_i \in \{0, 1\}$, $\mathbf{x}_i = (x_{i0}, x_{i1}, \dots, x_{ip})^\top$.
- $y_i \mid \mathbf{x}_i \sim \text{Bern}(p_i(\beta))$

$$- p_i(\beta) = \Pr(Y = 1 \mid \mathbf{X} = \mathbf{x}_i; \beta) = \left(1 + e^{-\beta^\top \mathbf{x}_i}\right)^{-1}$$

$$- \beta^T \mathbf{x}_i = \mathbf{x}_i^T \beta = \beta_0 + \sum_{j=1}^p x_{ij} \beta_j$$

- Bernoulli Likelihood Function

$$L(\beta) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

$$\log L(\beta) = \sum_{i=1}^n \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\}$$

- The usual approach to estimating the Logistic Regression coefficients is *maximum likelihood*

$$\begin{aligned} \hat{\beta} &= \arg \max_{\beta} L(\beta) \\ &= \arg \max_{\beta} \log L(\beta) \end{aligned}$$

- We can also view this as the coefficients that minimize the *loss function*, where the loss function is the negative log-likelihood

$$\begin{aligned} \hat{\beta} &= \arg \min_{\beta} \ell(\beta) \\ &= -C \sum_{i=1}^n \{y_i \ln p_i + (1 - y_i) \ln(1 - p_i)\} \end{aligned}$$

- where C is some constant, e.g., $C = 1/n$

- This view facilitates *penalized logistic regression*

$$\hat{\beta} = \arg \min_{\beta} \ell(\beta) + \lambda P(\beta)$$

- Ridge Penalty

$$P(\beta) = \sum_{j=1}^p |\beta_j|^2 = \beta^T \beta$$

- Lasso Penalty

$$P(\beta) = \sum_{j=1}^p |\beta_j|$$

- Best Subsets

$$P(\beta) = \sum_{j=1}^p |\beta_j|^0 = \sum_{j=1}^p 1_{(\beta_j \neq 0)}$$

3.3 Logistic Regression in Action

- In **R**, logistic regression can be implemented with the `glm()` function since it is a type of *Generalized Linear Model*.
- Because logistic regression is a special case of *Binomial* regression, use the `family=binomial()` argument

```
##-- Fit logistic regression model
fit.lr = glm(y~student + balance + income, data=Default,
            family="binomial")
```

term	estimate	std.error	statistic	p.value
(Intercept)	-10.8690	0.4923	-22.0801	0.0000
studentYes	-0.6468	0.2363	-2.7376	0.0062
balance	0.0057	0.0002	24.7376	0.0000
income	0.0000	0.0000	0.3698	0.7115

Your Turn #4 : Interpreting Logistic Regression

1. What is the estimated probability of default for a Student with a balance of \$1000?
2. What is the estimated probability of default for a *Non-Student* with a balance of \$1000?
3. Why does `student=Yes` appear to lower risk of default, when the plot of student status vs. default appears to increase risk?

4 Linear/Quadratic Discriminant Analysis (LDA/QDA)

- Discriminant analysis seeks to find a function that will *discriminate* between class boundaries.
 - Linear Discriminant Analysis (LDA)* finds *linear* boundaries between classes
 - Quadratic Discriminant Analysis (QDA)* finds *quadratic* boundaries between classes

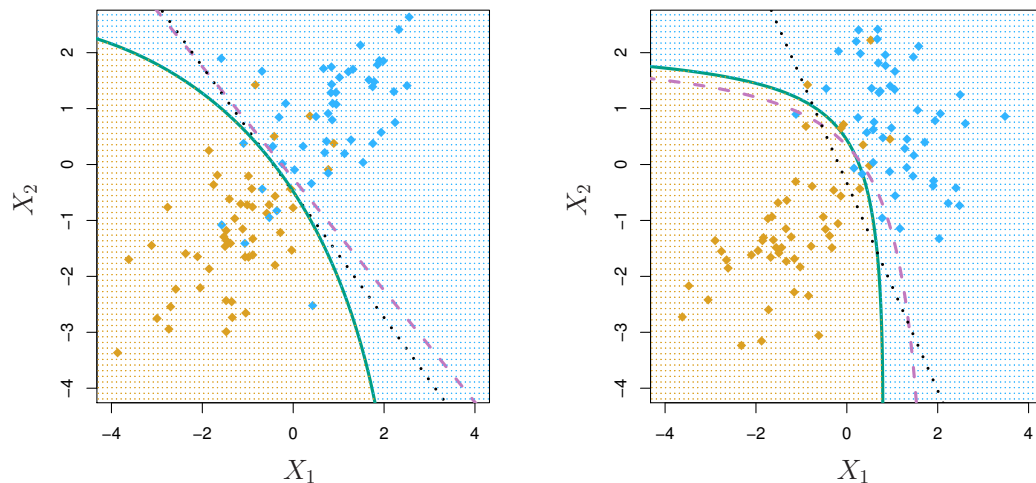


FIGURE 4.9 (from ISLR). Left: The Bayes (purple dashed), LDA (black dotted), and QDA (green solid) decision boundaries for a two-class problem with $\Sigma_1 = \Sigma_2$. The shading indicates the QDA decision rule. Since the Bayes decision boundary is linear, it is more accurately approximated by LDA than by QDA. Right: Details are as given in the left-hand panel, except that $\Sigma_1 \neq \Sigma_2$. Since the Bayes decision boundary is non-linear, it is more accurately approximated by QDA than by LDA.

- Suppose there are $K = |\mathcal{G}|$ classes in the training data, $D = \{(\mathbf{X}_i, G_i)\}_{i=1}^n$
 - where $\mathbf{X}_i \in \mathbf{R}^p$, $G_i \in \mathcal{G}$
- Consider the posterior probability of class g , given $X = x$,

$$\begin{aligned} \Pr(G = g \mid \mathbf{X} = \mathbf{x}) &= \frac{f(x \mid G = g) \Pr(G = g)}{f(x)} \\ &= \frac{f_g(x) \pi_g}{\sum_{k=1}^K f_k(x) \pi_k} \end{aligned}$$

- $f_k(x)$ is the *class conditional density*
 - $0 \leq \pi_k \leq 1$ are the *prior class probabilities*
 - $\sum_{k=1}^K \pi_k = 1$
- The challenge is to estimate the densities $\{f_k(\cdot)\}$
 - Note: $\hat{\pi}_k = n_k/n$ is a natural estimate for the class priors if we think the testing data will have the same proportions as the training data

4.1 Estimation

- LDA and QDA have strong connections to model based clustering.
 - But easier, since we have the true class labels in the supervised setting
- Both LDA and QDA model the class conditional densities $f_k(x)$ with *Gaussians*
 - Thus, they model the observations as coming from a *Gaussian mixture model*
 - Each class has its own mean vector μ_k
 - The difference between LDA and QDA is what they use for their covariance matrix

- **LDA**

$$f_k(x) = (2\pi)^{-p/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu_k)^\top \Sigma^{-1} (\mathbf{x} - \mu_k) \right\}$$

- $\Sigma_k = \Sigma \quad \forall k$ (uses the same variance-covariance for all classes)

- **QDA**

$$f_k(x) = (2\pi)^{-p/2} |\Sigma_k|^{-1/2} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \mu_k)^\top \Sigma_k^{-1} (\mathbf{x} - \mu_k) \right\}$$

- Σ_k is different for each classes

Your Turn #5 : Model Complexity

The LDA model uses a common covariance matrix while QDA allows each class to have a different covariance (which permits quadratic boundaries). But this flexibility comes at a cost.

1. How many parameters have to be estimated in an LDA model with K classes and p dimensions?
2. How many parameters have to be estimated in an QDA model with K classes and p dimensions?

- There are a few methods to maintain some flexibility, yet protect the model from high variance
- One is to use a *regularized covariance matrix* (see ESL 4.3.1)

$$\hat{\Sigma}_k(\alpha, \gamma) = \alpha \hat{\Sigma}_k + (1 - \alpha) \{ \gamma \hat{\Sigma} + (1 - \gamma) \hat{\sigma}^2 I_p \}$$

- Another is to fit an LDA model in an *enlarged feature space*
 - E.g., for $p = 2$ dimensions, use $X_1, X_2, X_1 \cdot X_2, X_1^2, X_2^2$ instead of QDA in X_1, X_2 .
 - Think basis expansion like what we say with polynomial regression or B-splines

4.2 LDA/QDA in Action

- In **R**, LDA and QDA can be implemented with the `lda()` and `qda()` functions from the **MASS** package.
- See ISLR 4.6 for details

5 Evaluating Classification Models

- Training Data: $\{X_i, G_i\}$
 - $G_i \in \{1, \dots, K\}$ (i.e., there are K classes)
- Predictor: $\hat{G}(X)$
- Loss function: $L(G, \hat{G}(X))$ is the loss incurred by estimating G with \hat{G}

- Risk is the expected loss (or expected prediction error EPE)
 - Expectation is taken wrt future values of (X, G)

$$\begin{aligned}
 \text{Risk}(\hat{G}) &= \text{EPE} \\
 &= E_{XG} [L(G, \hat{G}(X))] \\
 &= E_X [E_{G|X} [L(G, \hat{G}(X)) | X]] \\
 &= E_X [R_X(\hat{G})]
 \end{aligned}$$

- The Risk at input $X = x$ is

$$\begin{aligned}
 R_x(\hat{G}) &= E_{G|X=x} [L(G, \hat{G}(x)) | X = x] \\
 &= \sum_{k=1}^K L(G = k, \hat{G}(x)) \Pr(G = k | X = x)
 \end{aligned}$$

- Thus the optimal class label, given $X = x$, is

$$\hat{G}(x) = \arg \min_g R_x(g)$$

5.1 Evaluation of Binary Classification Models

- We are considering *binary* outcomes, so use the notation $Y \in \{0, 1\}$
- Let $p(x) = \Pr(Y = 1 | X = x)$
- The Risk (for a binary outcome) is:

$$\begin{aligned}
 R_x(g) &= L(1, g) \Pr(Y = 1 | X = x) + L(0, g)(1 - \Pr(Y = 1 | X = x)) \\
 &= L(1, g)p(x) + L(0, g)(1 - p(x))
 \end{aligned}$$

- Decision: choose $\hat{G}(x) = 1$ if

$$\begin{aligned}
 R_x(1) &< R_x(0) \\
 L(1, 1)p(x) + L(0, 1)(1 - p(x)) &< L(1, 0)p(x) + L(0, 0)(1 - p(x)) \\
 p(x)(L(1, 1) - L(1, 0)) &< (1 - p(x))(L(0, 0) - L(0, 1)) \\
 p(x)(L(1, 0) - L(1, 1)) &\geq (1 - p(x))(L(0, 1) - L(0, 0)) \quad (\text{multiply both sides by } -1) \\
 \frac{p(x)}{1 - p(x)} &\geq \frac{L(0, 1) - L(0, 0)}{L(1, 0) - L(1, 1)}
 \end{aligned}$$

5.1.1 Example

- Say we have a goal of estimating if a patient has cancer using medical imaging
 - Let $G = 1$ for cancer and $G = 0$ for no cancer
- Suppose we have solicited a loss function with the following values
 - $L(G = 0, \hat{G} = 0) = 0$: There is no loss for correctly diagnosis a patient without cancer
 - $L(G = 1, \hat{G} = 1) = 0$: There is no loss (for our model) for correctly diagnosis a patient with cancer
 - $L(G = 0, \hat{G} = 1) = FP$: There is a cost of FP units if the model issues a *false positive*, estimating the patient has cancer when they don't
 - $L(G = 1, \hat{G} = 0) = FN$: There is a cost of FN units if the model issues a *false negative*, estimating the patient does not have cancer when they really do

- In these scenarios FN is often much larger than FP ($FN \gg FP$) because the effects of not promptly treating (or further testing, etc) a patient is more severe than starting a treatment path for patients that don't actually have cancer
- Our model will decide to issue a positive indication for cancer if $R_x(1) < R_x(0)$ which occurs when

$$\frac{p(x)}{1 - p(x)} \geq \frac{FP}{FN}$$

$$p(x) \geq \frac{FP}{FP + FN}$$

- The ratio of FP to FN is all that matters for the decision. Let's say that $FP=1$ and $FN=10$. Then if $p(x) \geq 1/11$, our model will diagnose cancer.
 - Note: $p(x) = \Pr(Y = 1 | X = x)$ is affected by the class prior $\Pr(Y = 1)$ (e.g., the portion of the population tested with cancer), which is usually going to be small.

5.2 Common Binary Loss Functions

- Suppose we are going to estimate a binary response $Y \in \{0, 1\}$ with a (possibly continuous) predictor $\hat{y}(x)$
- **0-1 Loss or Misclassification Error**

$$L(y, \hat{y}(x)) = \mathbb{1}(y \neq \hat{y}(x)) = \begin{cases} 0 & y = \hat{y}(x) \\ 1 & y \neq \hat{y}(x) \end{cases}$$

- This assumes $L(0, 1) = L(1, 0)$ (i.e., false positive costs the same as a false negative)
- Requires that a *hard classification* is made
- The optimal prediction is $y^*(x) = \mathbb{1}(p(x) > 1 - p(x))$ which is $\mathbb{1}(p(x) > 0.50)$
- **Squared Error**

$$L(y, \hat{y}(x)) = (y - \hat{y}(x))^2$$

- The optimal prediction is $y^*(x) = E[Y | X = x] = \Pr(Y = 1 | X)$
- **Absolute Error**

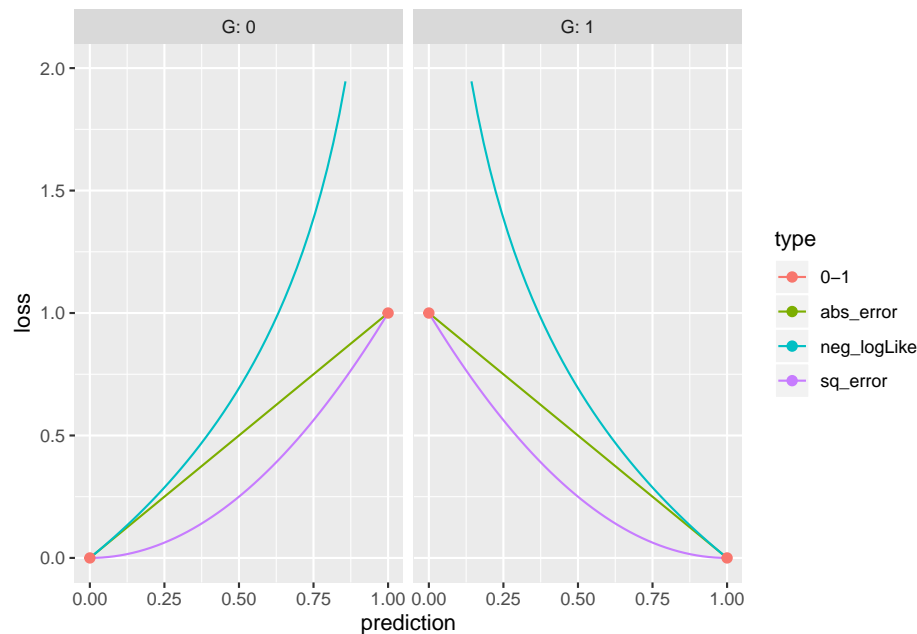
$$L(y, \hat{y}(x)) = |y - \hat{y}(x)|$$

- **Bernoulli negative log-likelihood**

$$L(y, \hat{y}(x)) = -\{y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)\}$$

$$= \begin{cases} -\log \hat{y} & y = 1 \\ -\log(1 - \hat{y}) & y = 0 \end{cases}$$

- Requires $\hat{y}(x) \in [0, 1]$



5.3 Evaluating Binary Classification Models

- Recall, the optimal (*hard classification*) decision is to choose $\hat{G} = 1$ if:

$$\frac{p(x)}{1 - p(x)} \geq \frac{L(0, 1) - L(0, 0)}{L(1, 0) - L(1, 1)}$$

- Denote $\gamma(x)$ as the *logit* of $p(x)$:

$$\gamma(x) = \log \frac{p(x)}{1 - p(x)} = \log \frac{\Pr(G = 1 \mid X = x)}{\Pr(G = 0 \mid X = x)}$$

- Then we get

$$p(x) = \Pr(G = 1 \mid X = x) = \frac{e^{\gamma(x)}}{1 + e^{\gamma(x)}}$$

- And the optimal (*hard classification*) decision can be described in terms of $\gamma(x)$:

Choose $\hat{G}(x) = 1$ if $\hat{\gamma}(x) > t$, where t is a threshold

- If the loss/cost is known, then

$$t^* = \log \left(\frac{L(0, 1) - L(0, 0)}{L(1, 0) - L(1, 1)} \right)$$

- For a given threshold t and input x , the hard classification is $\hat{G}_t(x) = \mathbb{1}(\hat{\gamma}(x) \geq t)$

5.4 Performance Metrics

5.4.1 Metrics

Metric	Definition	Estimate
Mis-classification Rate	$P_{XG}(\hat{G}_t(X) \neq G(X)) =$ $P_X(\hat{G}_t(X) = 0, G(X) = 1) +$ $P_X(\hat{G}_t(X) = 1, G(X) = 0)$	$\frac{1}{N} \sum_{i=1}^N \mathbb{1}(\hat{G}_t(x_i) \neq G_i)$
False Positive Rate (FPR) {1-Specificity}	$P_X(\hat{G}_t(X) = 1 \mid G(X) = 0)$	$\frac{1}{N_0} \sum_{i:G_i=0} \mathbb{1}(\hat{G}_t(x_i) = 1)$
True Positive Rate (TPR) {Hit Rate, Recall, Sensitivity}	$P_X(\hat{G}_t(X) = 1 \mid G(X) = 1)$	$\frac{1}{N_1} \sum_{i:G_i=1} \mathbb{1}(\hat{G}_t(x_i) = 1)$
Precision	$P_X(G(X) = 1 \mid \hat{G}_t(X) = 1)$	$\frac{1}{\hat{N}_1(t)} \sum_{i:\hat{G}_t(x_i)=1} \mathbb{1}(G_i = 1)$

- Note: Performance estimates are best carried out on *hold-out* data!

5.4.2 Confusion Matrix

- Given a threshold t , we can make a *confusion matrix* to help analyze our model's performance on data
 - Data = $\{(X_i, G_i)\}_{i=1}^N$ (ideally this is hold-out/test data)
 - N_k is number of observations from class k ($N_0 + N_1 = N$)

		Model Outcome		total
		$\hat{G}_t = 1$	$\hat{G}_t = 0$	
True Outcome	$G = 1$	True Positive (TP)	False Negative (FN)	N_1
	$G = 0$	False Positive (FP)	True Negative (TN)	N_0
total		$\hat{N}_1(t)$	$\hat{N}_0(t)$	N

Table from: <https://tex.stackexchange.com/questions/20267/how-to-construct-a-confusion-matrix-in-latex>

- See [Wikipedia Page: Confusion Matrix](#) for more metrics

5.5 Performance over a range of thresholds

To illustrate how we can calculate and visualize different performance aspects in binary classification let's go back to the `Default` data and see how well the basic logistic regression models works.

- In order to evaluation on hold-out data, split the data into train/test (used about 75% training, 25% testing), fit a logistic regression model on training data, and make predictions on the test data


```

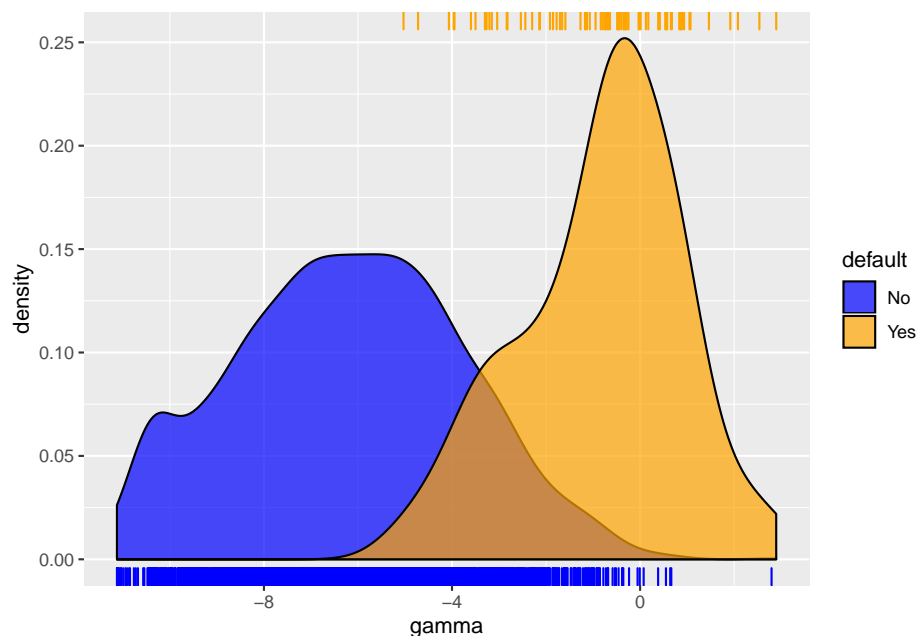
#-- train/test split
set.seed(2019)
Default = Default %>%
  mutate(group = sample(c('train', 'test'), size=nrow(.),
                        replace=TRUE, prob=c(.75, .25) )) # ~75% train

#-- fit model on training data
fit.lm = glm(y~student + balance + income,
             family='binomial',
             data=filter(Default, group=='train'))

#-- Get predictions (of gamma(x)) on test data
gamma = predict(fit.lm,
               newdata=filter(Default, group=='test'),
               type='link')
Gtest = filter(Default, group=='test') %>% pull(y) # true values
#> Error in path.expand(path): invalid 'path' argument

```

- The model is unable to perfectly discriminate between groups, but the *defaults* do get scored higher in general:
 - As a reference point, note that $\hat{\gamma}(x) = 0 \rightarrow \Pr(Y = 1 \mid X = x) = 1/2$



- We can calculate performance over a range of thresholds.
 - Unless the test data is too large, use all unique values of the training data as the thresholds

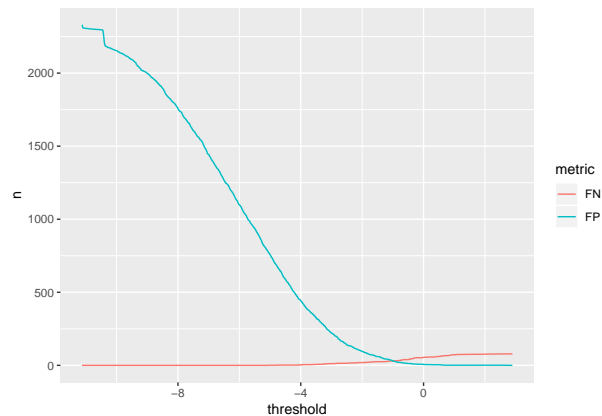
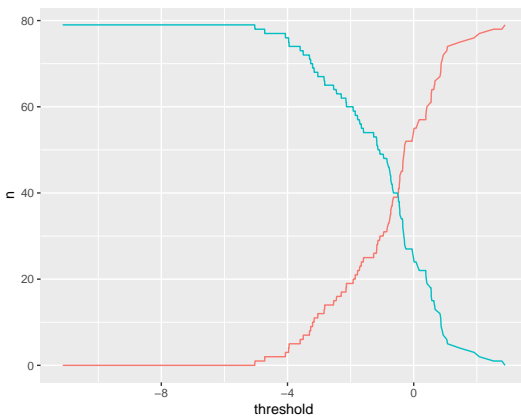
```

#-- Get performance data (by threshold)
perf = tibble(truth = Gtest, prediction = gamma) %>%
  #- group_by() + summarize() in case of ties
  group_by(prediction) %>%
  summarize(n=n(), n.1=sum(truth), n.0=n-sum(truth)) %>%
  #- calculate metrics
  arrange(prediction) %>%

```

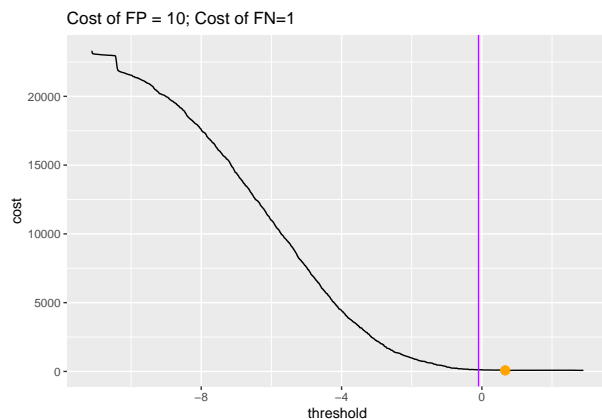
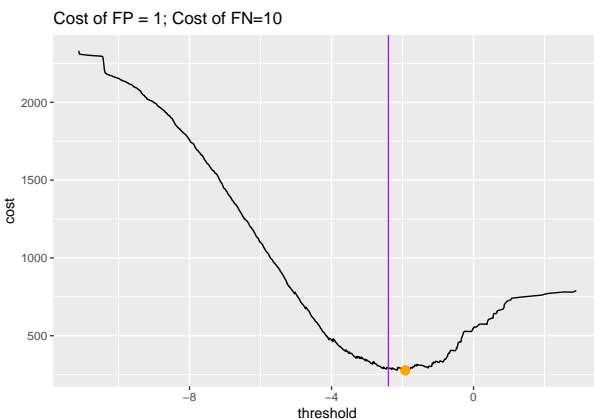
```
mutate(FN = cumsum(n.1),      # false negatives
       TN = cumsum(n.0),      # true positives
       TP = sum(n.1) - FN,     # true positives
       FP = sum(n.0) - TN,     # false positives
       N = cumsum(n),          # number of cases predicted to be 1
       TPR = TP/sum(n.1), FPR = FP/sum(n.0)) %>%
  #- only keep relevant metrics
  select(-n, -n.1, -n.0, threshold=prediction)
```

- **General Performance** (select metrics)



- **Cost Curves**

- note: the **purple** is the *theoretical* optimal threshold (using $t^* = \log FP/FN$) and the **orange** point is at the optimal value using the model

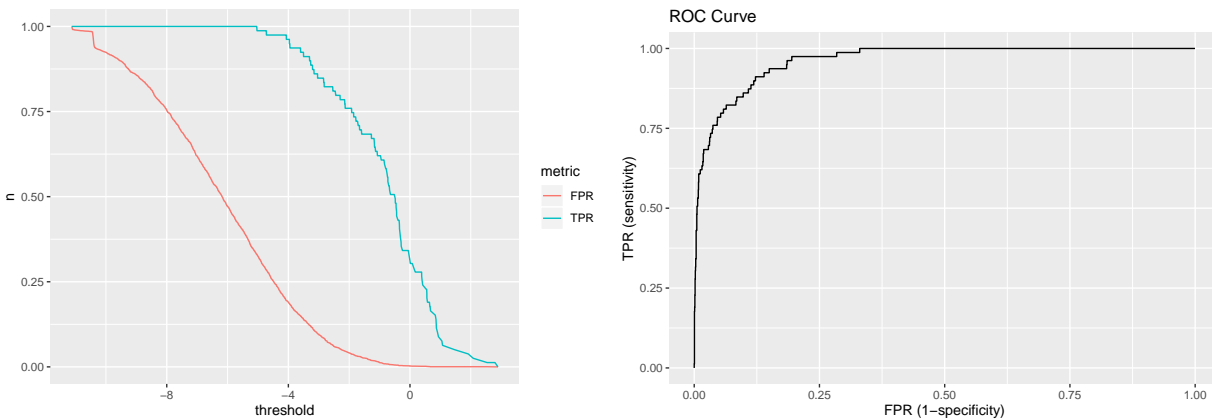


Optimal Threshold

- The *theoretically* optimal threshold is based on the *true* $\gamma(x) = \log \frac{p(x)}{1-p(x)}$ (for a given cost ratio of FP to FN)
- The observed optimal threshold will differ when the model's estimate $\hat{\gamma}(x) \neq \gamma(x)$
 - Hopefully, they are close and it won't make much difference which one you use. But I'd take the estimated threshold if I had sufficient data.
- Note that the estimated values depend on the prior class probabilities. If you suspect these may differ in

the future, then you should adjust the threshold.

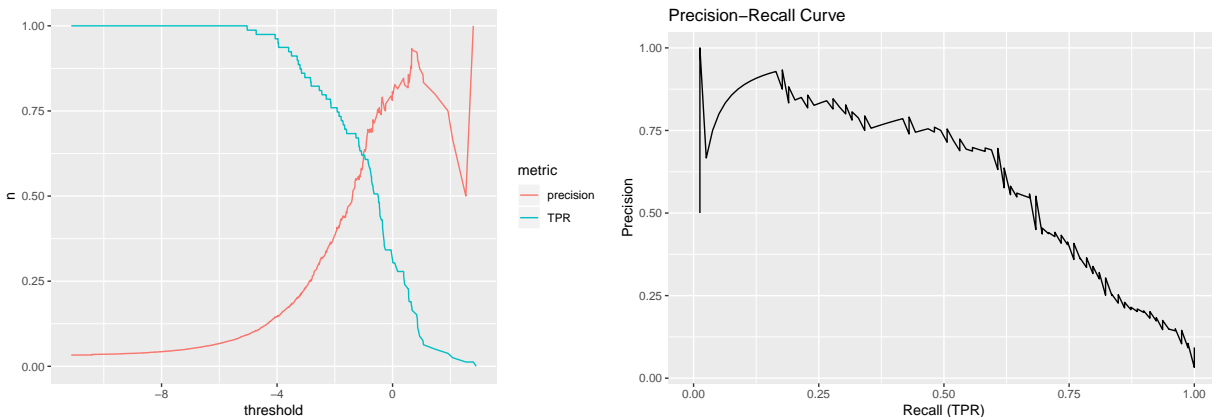
- **ROC Curves** (Receiver Operating Characteristic)



- The *area under the ROC curve* (AUC) is a popular performance metric
- I don't think it is a great way to compare classifiers for several reasons
 - The main reason is that in a real application you can almost always come up with an estimated cost/loss for the different decisions
 - To say it another way, comparisons should be made at a single point on the curve; the entire FPR region should not factor into the comparison.

- **Precision Recall Curves**

- Popular for information retrieval/ranking



6 Naive Bayes and Generalized Additive Models (GAM)

6.1 The Bayes Breakdown

Recall our notation for binary classification problems

- $p(x) = \Pr(Y = 1 \mid X = x) = \frac{f_1(x)\pi}{f_1(x)\pi + f_0(x)(1-\pi)}$
- $\gamma(x) = \frac{p(x)}{1-p(x)}$

- $f_k(x)$ is the *class conditional density*
- $0 \leq \pi_k \leq 1$ are the *prior class probabilities*
- $\pi_0 + \pi_1 = 1$

The log-odds reduces to a combination of prior odds and density (likelihood) ratios

$$\begin{aligned}\gamma(x) &= \log \left(\frac{p(x)}{1 - p(x)} \right) \\ &= \underbrace{\log \left(\frac{\pi}{1 - \pi} \right)}_{\text{log prior odds}} + \underbrace{\log \left(\frac{f_1(x)}{f_0(x)} \right)}_{\text{log density ratio}}\end{aligned}$$

- Note: $\frac{f_1(x)}{f_0(x)}$ is usually called a *likelihood ratio* when estimated via MLE and *Bayes Factor* when integrating over the model parameters
- We can see that the optimal decision can be based on the density ratios

Choose $\hat{G}(x) = 1$ if:

$$\begin{aligned}\hat{\gamma}(x) &> \log \left(\frac{L(0, 1) - L(0, 0)}{L(1, 0) - L(1, 1)} \right) \\ &= \log \left(\widehat{\frac{f_1(x)}{f_0(x)}} \right) > \log \left(\frac{1 - \hat{\pi}}{\hat{\pi}} \right) + \log \left(\frac{L(0, 1) - L(0, 0)}{L(1, 0) - L(1, 1)} \right)\end{aligned}$$

6.2 Naive Bayes

- Recall in LDA/QDA, the class conditional densities were estimated as Gaussians:

$$\hat{f}_k(\mathbf{x}) = \mathcal{N}(\mathbf{x}; \hat{\mu}_k, \hat{\Sigma}_k)$$

- But when the dimensionality of \mathbf{x} gets large or there is high correlation, estimation of $\hat{\Sigma}_k$ can be poor
- Also if the multivariate densities are not unimodal these models may not have good predictive ability
- If we force $\hat{\Sigma}_k$ to be *diagonal* then the densities are product of univariate Gaussians

$$\hat{f}_k = \prod_{j=1}^p \mathcal{N}(x_j; \mu_{kj}, \sigma_{kj})$$

- Even if the data are not independent, this may give better estimates by reducing the variance (at the expense of a bit of bias)
- This is an example of *Naive Bayes*
- The so called **Naive Bayes** approach just models the joint multivariate densities as the product of marginal densities

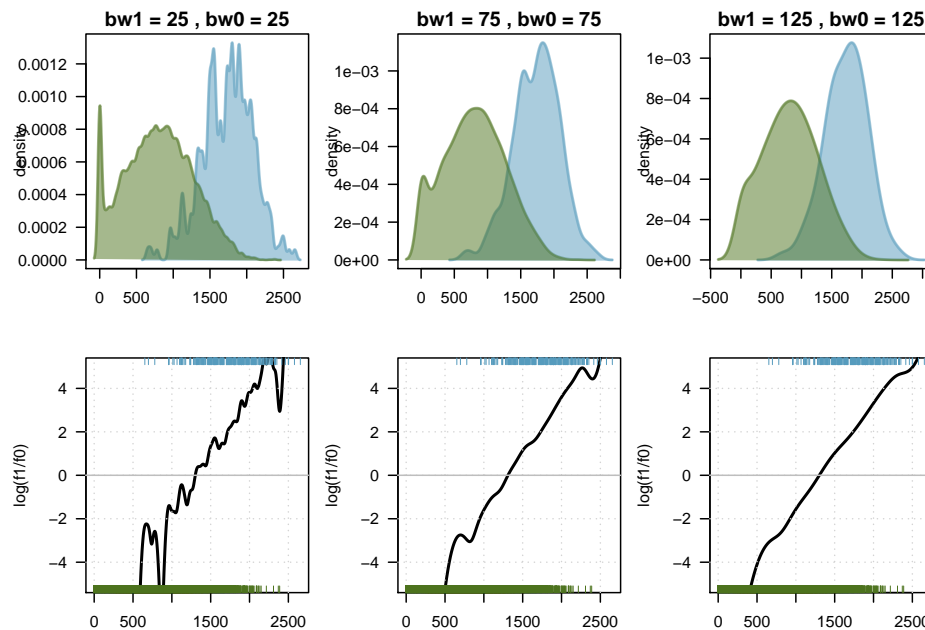
$$\hat{f}_k = \prod_{j=1}^p \hat{f}_{kj}(x_j)$$

- The densities do *not* have to be Gaussian
- Categorical densities (i.e., pmfs) can be thrown in the mix without a problem
- Because of the independence, this is easy to implement in parallel (and thus can be fast)

6.2.1 KDE based Naive Bayes

- Kernel density estimation (KDE) can be a great way to estimate the component densities
- The bandwidth tuning parameter allows a flexible shape
 - Using the balance predictor in the Default data

```
#> Error in path.expand(path): invalid 'path' argument
#> Error in path.expand(path): invalid 'path' argument
```



6.3 Logistic Regression vs. Linear Discriminant Analysis (LDA) vs. Naive Bayes

It turns out that there is a close connection between Logistic Regression, Naive Bayes, and LDA. To help see this, notice that all three methods can be written:

$$\begin{aligned}\gamma(x) &= \log\left(\frac{\pi}{1-\pi}\right) + \log\left(\frac{f_1(x)}{f_0(x)}\right) \\ &= \alpha_0 + \sum_{j=1}^p \alpha_j S_j\end{aligned}$$

- **Logistic Regression**

$$\hat{\alpha}_0 = \hat{\beta}_0$$

$$\hat{\alpha}_j = \hat{\beta}_j$$

$$\hat{S}_j = x_j$$

- **LDA**

$$\hat{\alpha}_0 = \log \frac{\hat{\pi}}{1-\hat{\pi}} - \frac{1}{2}(\hat{\mu}_1 + \hat{\mu}_0)^T \hat{\Sigma}^{-1}(\hat{\mu}_1 - \hat{\mu}_0)$$

$$\hat{\alpha}_j = \hat{\Sigma}^{-1}(\hat{\mu}_1 - \hat{\mu}_0)$$

$$\hat{S}_j = x_j$$

- **Naive Bayes**

$$\hat{\alpha}_0 = \log \frac{\hat{\pi}}{1 - \hat{\pi}}$$

$$\hat{\alpha}_j = 1$$

$$\hat{S}_j = \log \frac{\hat{f}_{1j}(x_j)}{\hat{f}_{0j}(x_j)}$$

- **Generalized Additive Models (GAM)**

- GAM models are made to directly estimate models of this form.

$$\hat{\gamma}(x) = \hat{\alpha} + \sum_{j=1}^p \hat{g}_j(x_j)$$

- In **R**, the `mgcv` package is worth becoming familiar with to implement GAM.
- See ESL 9.1 for more details