

R Formula Interface

and Design Matrices

SYS 4582/6018 | Spring 2019

Rfm1a.pdf

1 Raw input data

The raw input data is often in the form of a data frame. For example,

```
#-- Raw Input Data
# cat is categorical with 3 levels: A,B,C
# num is numerical
# y is numerical response variable

Z = data.frame(cat=c('A', 'A', 'B', 'B', 'C', 'C'), num=1:6, y=rnorm(6))
Z

#>   cat num      y
#> 1   A   1 -0.98753
#> 2   A   2  1.59598
#> 3   B   3  0.96065
#> 4   B   4 -0.03212
#> 5   C   5 -0.39614
#> 6   C   6 -0.95628
```

has three columns, `cat` is categorical data, `num` which is numerical data, and `y` which is the response variable.

2 Formula in models

The formula interface in R allows you to make transformations of the input data frame automatically. For example, categorical (or factor) columns will generate the appropriate dummy variables.

```
lm(y~cat, data=Z)$coef

#> (Intercept)      catB      catC
#>    0.3042    0.1600   -0.9804

lm(y~cat - 1, data=Z)$coef # remove intercept

#>   catA   catB   catC
#> 0.3042 0.4643 -0.6762
```

The default behavior is to convert categorical data to a *factor* and drop the first level.

The formula interface is easy to use:

```
#- numerical data only
lm(y~num, data=Z)$coef

#> (Intercept)      num
#>    0.7120   -0.1947
```

```

#- transformations
lm(y~log(num), data=Z)$coef

#> (Intercept)      log(num)
#>      0.2486      -0.1987

#- use I() to make custom functions
lm(y~I(3*num), data=Z)$coef

#> (Intercept)      I(3 * num)
#>      0.71205      -0.06488

#- we have already seen poly()
lm(y~poly(num, degree = 3), data=Z)$coef

#> (Intercept) poly(num, degree = 3)1 poly(num, degree = 3)2
#>      0.03076      -0.81429      -1.59659
#> poly(num, degree = 3)3
#>      1.34702

#- how about B-splines
library(splines)
lm(y~bs(num), data=Z)$coef

#> (Intercept)      bs(num)1      bs(num)2      bs(num)3
#>      -0.85563      5.67394      -1.28813      0.03075

#- two predictors
lm(y~cat + num, data=Z)$coef

#> (Intercept)      catB      catC      num
#>      -0.2111      -0.5270      -2.3546      0.3435

lm(y~cat + num - 1, data=Z)$coef

#>      catA      catB      catC      num
#> -0.2111 -0.7381 -2.5657  0.3435

#- a:b stands for interactions
lm(y~cat + num + cat:num, data=Z)$coef

#> (Intercept)      catB      catC      num      catB:num      catC:num
#>      -3.571      7.510      5.976      2.584      -3.576      -3.144

#- use . to represent everything in data
lm(y~., data=Z)$coef

#> (Intercept)      catB      catC      num
#>      -0.2111      -0.5270      -2.3546      0.3435

lm(y~. - num, data=Z)$coef # use . to include all, then remove some

#> (Intercept)      catB      catC
#>      0.3042      0.1600      -0.9804

```

2.1 model.matrix()

Behind the scenes, `lm()` is calling the function `model.matrix()` to construct the design matrix, or the real valued X matrix used for calculating the coefficients. You have to pass a formula object into `model.matrix()`.

```
fmla = formula(y~num+cat)
model.matrix(fmla, data=Z)

#>      (Intercept) num catB catC
#> 1             1   1    0    0
#> 2             1   2    0    0
#> 3             1   3    1    0
#> 4             1   4    1    0
#> 5             1   5    0    1
#> 6             1   6    0    1
#> attr(,"assign")
#> [1] 0 1 2 2
#> attr(,"contrasts")
#> attr(,"contrasts")$cat
#> [1] "contr.treatment"

fmla = formula(y~num+cat-1) # remove intercept
model.matrix(fmla, data=Z)
```

```
#>      num catA catB catC
#> 1     1     1     0     0
#> 2     2     1     0     0
#> 3     3     0     1     0
#> 4     4     0     1     0
#> 5     5     0     0     1
#> 6     6     0     0     1
#> attr(,"assign")
#> [1] 1 2 2 2
#> attr(,"contrasts")
#> attr(,"contrasts")$cat
#> [1] "contr.treatment"
```

Or, if you are good with data manipulation construct the design matrix manually.

```
library(dplyr)
transmute(Z, intercept=1,
          x1=num, x2=num^2,
          x3=ifelse(cat=='B',1,0), x4=ifelse(cat=='C',1,0)) %>% as.matrix

#>      intercept x1 x2 x3 x4
#> [1,]          1  1  1  0  0
#> [2,]          1  2  4  0  0
#> [3,]          1  3  9  1  0
#> [4,]          1  4 16  1  0
#> [5,]          1  5 25  0  1
#> [6,]          1  6 36  0  1
```

Some functions (e.g., `glmnet`) do not take formulas so you will have to pass in the design matrix X directly. Another word of caution, some functions (again like `glmnet`) add the intercept automatically so you should not include a columns of ones.

The function `lm.fit()` fits a linear model from a design matrix:

```
X = model.matrix(formula(y~num+cat), data=Z)
Y = Z$y
lm.fit(x=X, y=Y)$coef
```

```
#> (Intercept)          num          catB          catC
#>      -0.2111      0.3435      -0.5270      -2.3546
```

2.2 Comparison

It is always good to compare the approaches just to make sure there are no mistakes.

```
fmla = formula(y~num+cat + I(num^2) + sqrt(num))

#- lm()
beta.lm = lm(fmla, data=Z)$coef

#- lm.fit()
X = model.matrix(fmla, data=Z)
beta.lmfit = lm.fit(X, Z$y)$coef

#- direct matrix operations
beta.eq = solve(t(X) %*% X) %*% t(X) %*% Z$y

#- output
data.frame(beta.lm, beta.lmfit, beta.eq) %>% knitr::kable()
```

	beta.lm	beta.lmfit	beta.eq
(Intercept)	-27.7623	-27.7623	-27.7623
num	-17.6293	-17.6293	-17.6293
catB	-0.4149	-0.4149	-0.4149
catC	0.1993	0.1993	0.1993
I(num^2)	0.7039	0.7039	0.7039
sqrt(num)	43.7002	43.7002	43.7002