

R Formula Interface

and Model/Design Matrices

SYS 6018 | Spring 2022

Rfmla.pdf

```
#-- Required Packags
library(splines)
library(tidyverse)
```

1 Raw input data

The raw input data is often in the form of a data frame (or tibble). For example,

```
#-- Raw Input Data
# cat is categorical with 3 levels: A,B,C
# num is numerical
# y is numerical response variable

Z = tibble(cat=c('A','A','B','B','C','C'), num=1:6, y=rnorm(6))
Z
#> # A tibble: 6 x 3
#>   cat      num      y
#>   <chr> <int>   <dbl>
#> 1 A         1 -1.77
#> 2 A         2 -0.0951
#> 3 B         3  1.26
#> 4 B         4 -0.595
#> 5 C         5  2.69
#> 6 C         6  0.295
```

has three columns, `cat` is categorical data, `num` which is numerical data, and `y` which is the outcome variable.

2 Formula in models

The formula interface in R allows you to make transformations of the input data frame automatically. For example, categorical (or factor) columns will generate the appropriate dummy variables.

```
lm(y~cat, data=Z)$coef
#> (Intercept)      catB      catC
#>   -0.932      1.263      2.425
lm(y~cat - 1, data=Z)$coef # remove intercept
#>   catA   catB   catC
#> -0.9320 0.3307 1.4932
```

The default behavior is to convert categorical data to a *factor* and drop the first level.

The formula interface is easy to use:

```
#- numerical data only
lm(y~num, data=Z)$coef
#> (Intercept)      num
#>   -1.3854    0.4808

#- transformations
lm(y~log(num), data=Z)$coef
#> (Intercept)   log(num)
#>   -1.398    1.546

#- use I() to make custom functions
lm(y~I(3*num), data=Z)$coef
#> (Intercept)  I(3 * num)
#>   -1.3854    0.1603

#- we have already seen poly()
lm(y~poly(num, degree = 3), data=Z)$coef
#> (Intercept) poly(num, degree = 3)1 poly(num, degree = 3)2
#>   0.2973      2.0112      -1.3761
#> poly(num, degree = 3)3
#>   -0.1330

#- how about B-splines
library(splines)
lm(y~bs(num), data=Z)$coef
#> (Intercept)  bs(num)1  bs(num)2  bs(num)3
#>   -1.606    2.301    3.757    2.305

#- two predictors
lm(y~cat + num, data=Z)$coef
#> (Intercept)  catB  catC  num
#>   0.3553    2.9791    5.8579   -0.8582
lm(y~cat + num - 1, data=Z)$coef
#>  catA  catB  catC  num
#>  0.3553  3.3344  6.2132 -0.8582

#- a:b stands for interactions
lm(y~cat + num + cat:num, data=Z)$coef
#> (Intercept)  catB  catC  num  catB:num  catC:num
#>   -3.443    10.253    18.119    1.674    -3.525    -4.071

#- use . to represent everything in data
lm(y~., data=Z)$coef
#> (Intercept)  catB  catC  num
#>   0.3553    2.9791    5.8579   -0.8582
lm(y~. - num, data=Z)$coef # use . to include all, then remove some
#> (Intercept)  catB  catC
#>   -0.932    1.263    2.425
```

2.1 model.matrix()

Behind the scenes, `lm()` is calling the function `model.matrix()` to construct the *model matrix* (also known as a *design matrix*). The model matrix is the real valued X matrix used for calculating the coefficients. You have to pass a formula object into `model.matrix()`.

```
fmla = formula(y~num+cat)
model.matrix(fmla, data=Z)
```

```
#>      (Intercept) num catB catC
#> 1             1  1    0    0
#> 2             1  2    0    0
#> 3             1  3    1    0
#> 4             1  4    1    0
#> 5             1  5    0    1
#> 6             1  6    0    1
#> attr(,"assign")
#> [1] 0 1 2 2
#> attr(,"contrasts")
#> attr(,"contrasts")$cat
#> [1] "contr.treatment"

fmla = formula(y~num+cat-1) # remove intercept
model.matrix(fmla, data=Z)
#>      num catA catB catC
#> 1     1     1     0     0
#> 2     2     1     0     0
#> 3     3     0     1     0
#> 4     4     0     1     0
#> 5     5     0     0     1
#> 6     6     0     0     1
#> attr(,"assign")
#> [1] 1 2 2 2
#> attr(,"contrasts")
#> attr(,"contrasts")$cat
#> [1] "contr.treatment"
```

Or, if you are good with data manipulation construct the model matrix manually.

```
library(dplyr)
Z %>%
  transmute(intercept=1,
            x1=num, x2=num^2,
            x3=ifelse(cat=='B',1,0), x4=ifelse(cat=='C',1,0)) %>%
  as.matrix()
#>      intercept x1 x2 x3 x4
#> [1,]         1  1  1  0  0
#> [2,]         1  2  4  0  0
#> [3,]         1  3  9  1  0
#> [4,]         1  4 16  1  0
#> [5,]         1  5 25  0  1
#> [6,]         1  6 36  0  1
```

Some functions (e.g., `glmnet`) do not take formulas so you will have to pass in the model matrix X directly. Another word of caution, some functions (again like `glmnet`) add the intercept automatically so you should not include a columns of ones.

The function `lm.fit()` fits a linear model from a model matrix:

```
X = model.matrix(formula(y~num+cat), data=Z)
Y = Z$y
lm.fit(x=X, y=Y)$coef
#>      (Intercept)      num      catB      catC
#>      0.3553      -0.8582      2.9791      5.8579
```

2.2 Comparison

It is always good to compare the approaches just to make sure there are no mistakes.

```
fmla = formula(y~num+cat + I(num^2) + sqrt(num))
```

```
#- lm()
```

```
beta.lm = lm(fmla, data=Z)$coef
```

```
#- lm.fit()
```

```
X = model.matrix(fmla, data=Z)
```

```
beta.lmfit = lm.fit(X, Z$y)$coef
```

```
#- direct matrix operations
```

```
beta.eq = solve(t(X) %*% X) %*% t(X) %*% Z$y
```

```
#- output
```

```
tibble(beta.lm, beta.lmfit, beta.eq)
```

beta.lm	beta.lmfit	beta.eq
-21.8559	-21.8559	-21.8559
-12.6989	-12.6989	-12.6989
2.1953	2.1953	2.1953
7.7561	7.7561	7.7561
0.3064	0.3064	0.3064
32.4796	32.4796	32.4796